

IA-32 Processor Architecture

A Programmer Perspective

Modes of Operation

- Protected mode
 - native mode (Windows, Linux)
- Real-address mode
 - native MS-DOS
- System management mode
 - power management, system security, diagnostics

Multitasking

- Operating System can run multiple programs (processes) at the same time.
- Multiple threads of execution within the same process.
- Scheduler utility assigns a given amount of CPU time to each running program.
- Rapid switching of tasks
 - gives illusion that all programs are running at once
 - the processor must support task switching.

Basic Execution Environment

- Addressable memory
- General-purpose registers
- Index and base registers
- Specialized register uses
- Status flags
- Floating-point, MMX, XMM registers

Addressable Memory

- Protected mode
 - 4 GB
 - 32-bit address
- Real-address and Virtual-8086 modes
 - 1 MB space
 - 20-bit address

General-Purpose Registers

Named storage locations inside the CPU, optimized for speed.

32-bit General-Purpose Registers

EAX
EBX
ECX
EDX

EBP
ESP
ESI
EDI

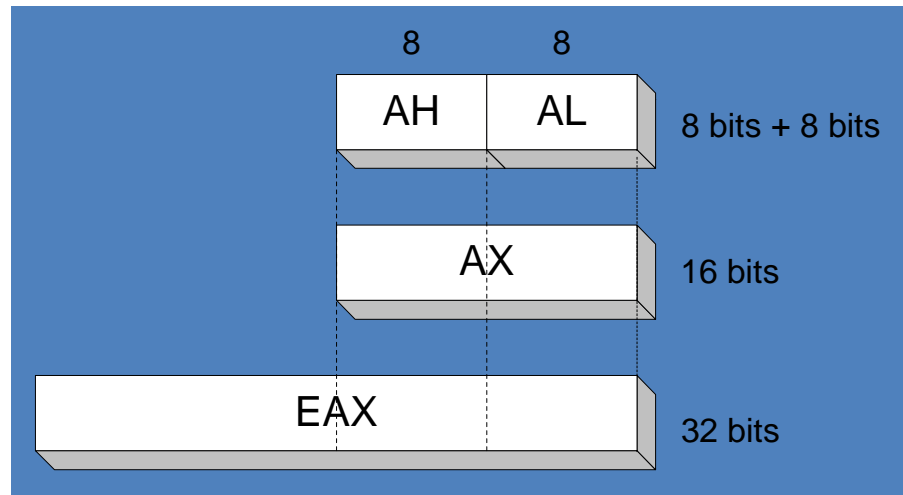
16-bit Segment Registers

EFLAGS
EIP

CS	ES
SS	FS
DS	GS

Accessing Parts of Registers

- Use 8-bit name, 16-bit name, or 32-bit name
- Applies to EAX, EBX, ECX, and EDX



32-bit	16-bit	8-bit (high)	8-bit (low)
EAX	AX	AH	AL
EBX	BX	BH	BL
ECX	CX	CH	CL
EDX	DX	DH	DL

Index and Base Registers

- Some registers have only a 16-bit name for their lower half
- The 16-bit registers are usually used only in real-address mode

32-bit	16-bit
ESI	SI
EDI	DI
EBP	BP
ESP	SP

Some specialized register uses

- General-Purpose
 - EAX – accumulator (automatically used by division and multiplication)
 - ECX – loop counter
 - ESP – stack pointer (should never be used for arithmetic or data transfer)
 - ESI, EDI – index registers (used for high-speed memory transfer instructions)
 - EBP – extended frame pointer (stack)

Some specialized register uses (cont.)

- Segment
 - CS – code segment
 - DS – data segment
 - SS – stack segment
 - ES, FS, GS - additional segments
- EIP – instruction pointer
- EFLAGS
 - control flags (control CPU's operation, e.g. break, interrupt, enter 8086/protected mode)
 - Status flag
 - each flag is a single binary bit (*set* or *clear*)

Status Flags

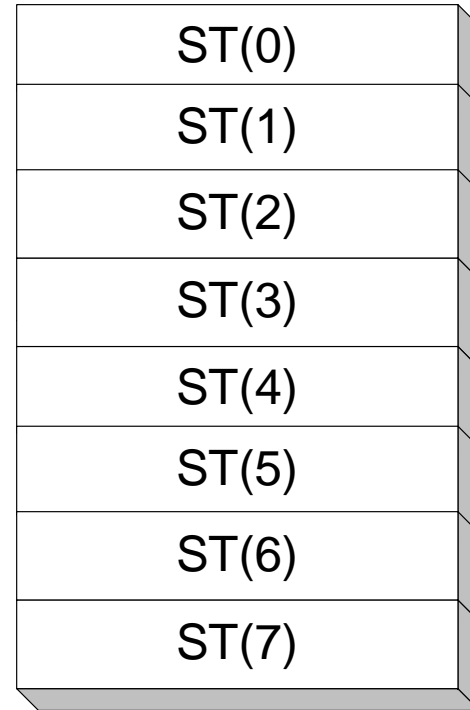
- Carry (CF)
 - unsigned arithmetic out of range
- Overflow (OF)
 - signed arithmetic out of range
- Sign (SF)
 - result is negative
- Zero (ZF)
 - result is zero
- Auxiliary Carry (AC)
 - carry from bit 3 to bit 4 in 8-bit operand
- Parity (PF)
 - sum of 1 bits in least-significant byte is an even number

System registers

- Accessed by operating system kernel at highest privilege level, not by application programs
 - IDTR (Interrupt Descriptor Table Register)
 - GDTR (Global Descriptor Table Register)
 - LDTR (Local Descriptor Table Register)
 - Task register
 - Debug registers
 - Control registers (e.g. task switching, paging, enabling cache memory)
 - Model-specific registers (e.g. performance monitoring, checking architecture)

Floating-Point, MMX, XMM Registers

- Eight 80-bit floating-point data registers
 - ST(0), ST(1), . . . , ST(7)
 - arranged in a stack
 - used for all floating-point arithmetic
- Eight 64-bit MMX registers
- Eight 128-bit XMM registers for single-instruction multiple-data (SIMD) operations



ST(0)
ST(1)
ST(2)
ST(3)
ST(4)
ST(5)
ST(6)
ST(7)

SIMD: A single computer instruction perform the same identical action (retrieve, calculate, or store) simultaneously on two or more pieces of data

IA-32 Memory Management

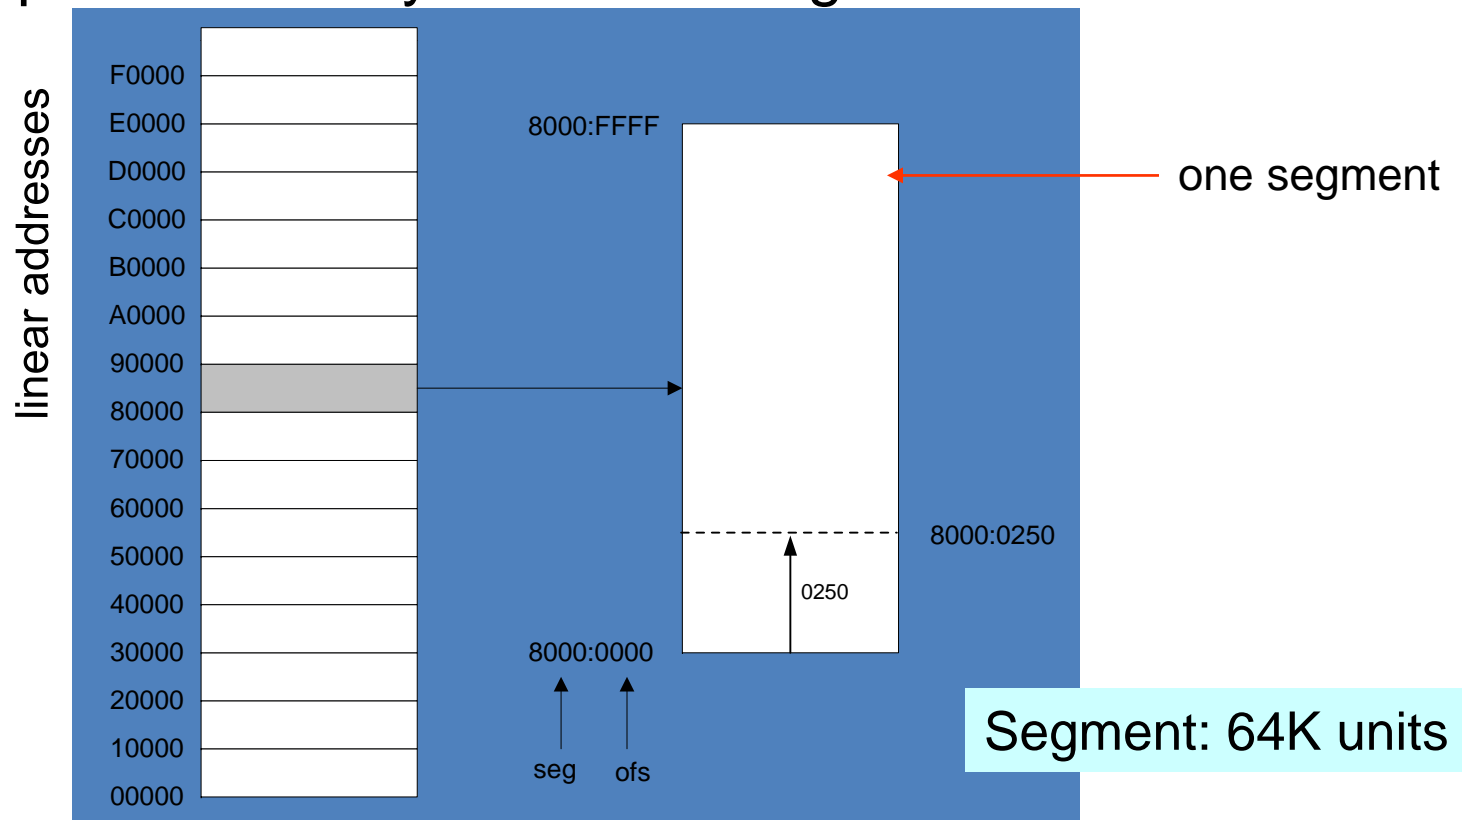
- Real-address mode
- Calculating linear addresses
- Protected mode
- Multi-segment model
- Paging

Real-Address mode

- 1 MB RAM maximum addressable
- Application programs can access any area of memory
- Single tasking
- Supported by MS-DOS operating system

Segmented Memory

- Segmented memory addressing: absolute (linear) address is a combination of a 16-bit segment value added to a 16-bit offset
- 8086 processor only has 16-bit registers



Calculating Linear Addresses

- Given a segment address, multiply it by 16 (add a hexadecimal zero), and add it to the offset
- Example: convert 08F1:0100 to a linear address

Adjusted Segment value:	0	8	F	1	0
Add the offset:			0	1	0 0
Linear address:			0	9	0 1 0

- A typical program has three segments: code, data and stack. Segment registers CS, DS and SS are used to store them separately.

Your turn . . .

What linear address corresponds to the segment/offset address 028F:0030?

$$028F0 + 0030 = 02920$$

Always use hexadecimal notation for addresses.

Your turn . . .

What segment addresses correspond to the linear address 28F30h?

Many different segment-offset addresses can produce the linear address 28F30h. For example:

28F0:0030, 28F3:0000, 28B0:0430, . . .

Protected Mode

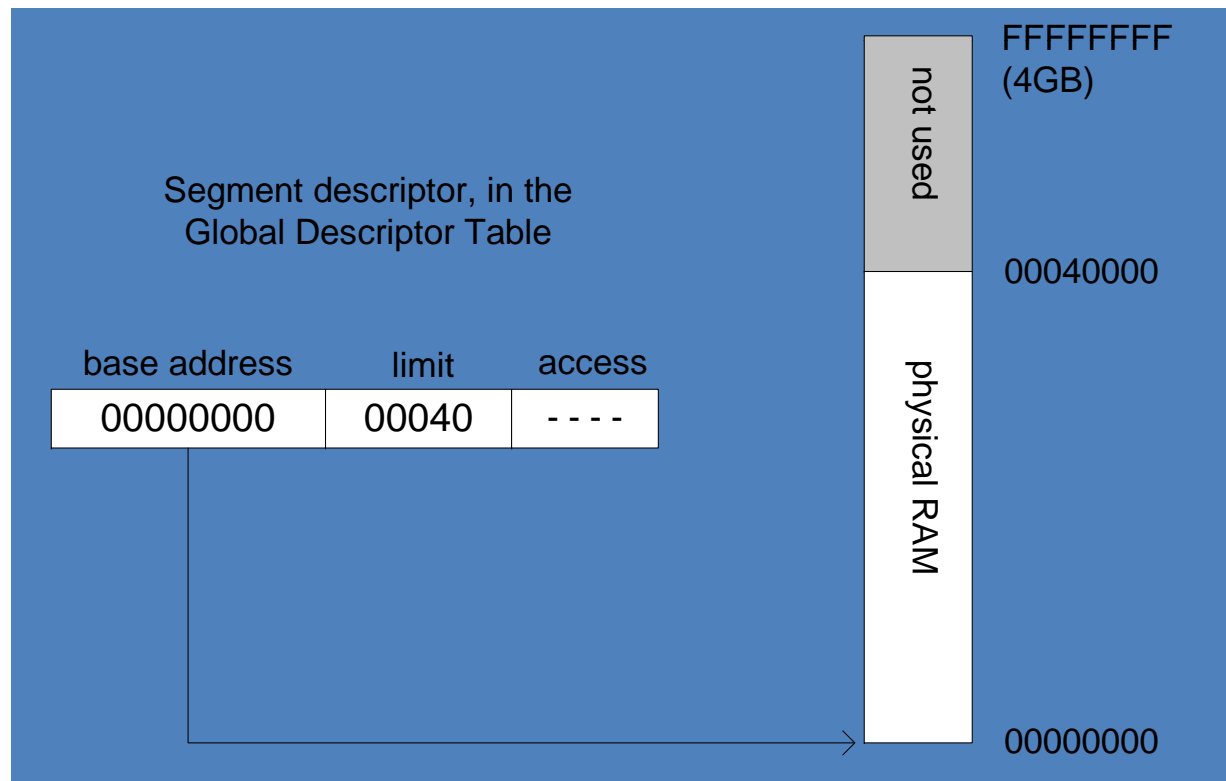
- 4 GB addressable RAM
 - (00000000 to FFFFFFFFh)
- Each program assigned a memory partition which is protected from other programs
- Designed for multitasking
- Supported by Linux & MS-Windows

Protected Mode (cont.)

- Segments
 - variable-sized areas of memory for code & data
- Segment descriptor
 - 64-bit value identifying and describing a single memory segment
 - contains segment's base address, access rights, size limit, type, and usage
- Segment descriptor table
 - contains segment descriptors by which OS keep track of locations of individual program segments
- Segment registers
 - points to segment descriptor tables
- Program structure
 - code, data, and stack areas
 - CS, DS, SS segment descriptors
 - global descriptor table (GDT)
- MASM Programs use the Microsoft **flat** memory model

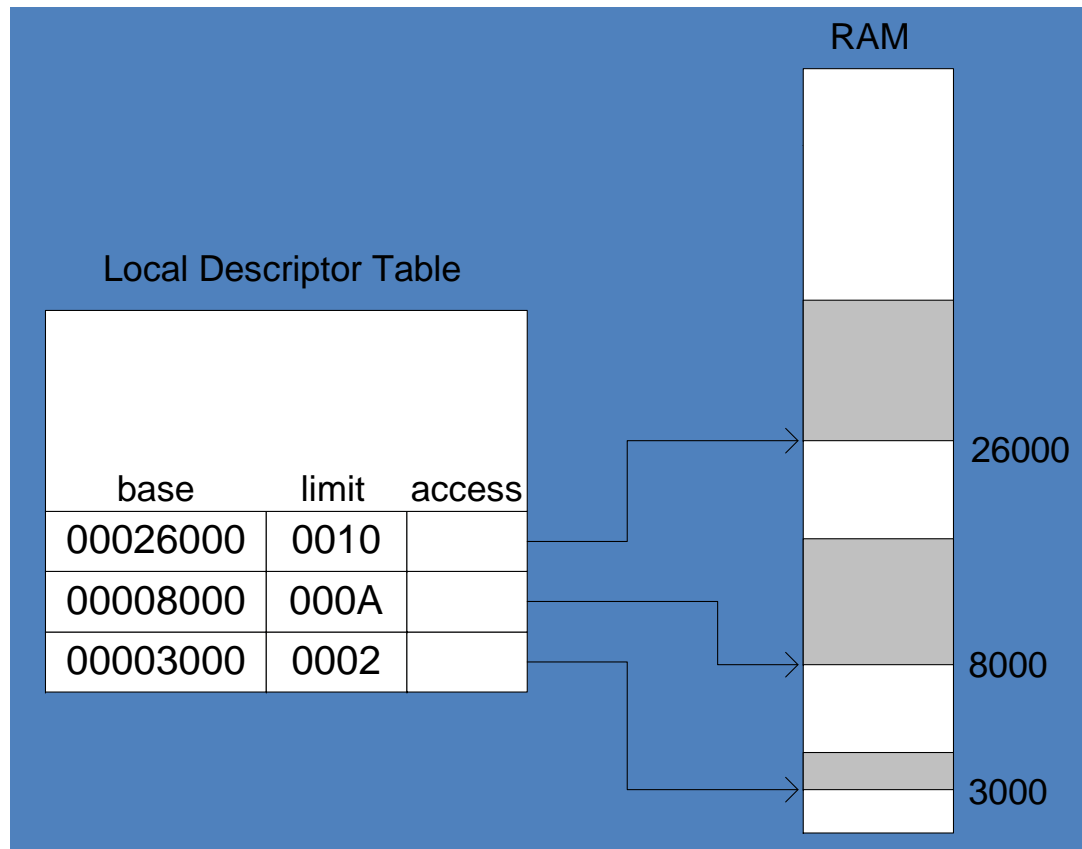
Flat Segment Model

- Single global descriptor table (GDT) whose base address is in GDTR
- Created when OS switches the processor into protected mode during boot up
- All segments mapped to entire 32-bit address space



Multi-Segment Model

- Each program has a local descriptor table (LDT)
 - holds descriptor for each segment used by the program

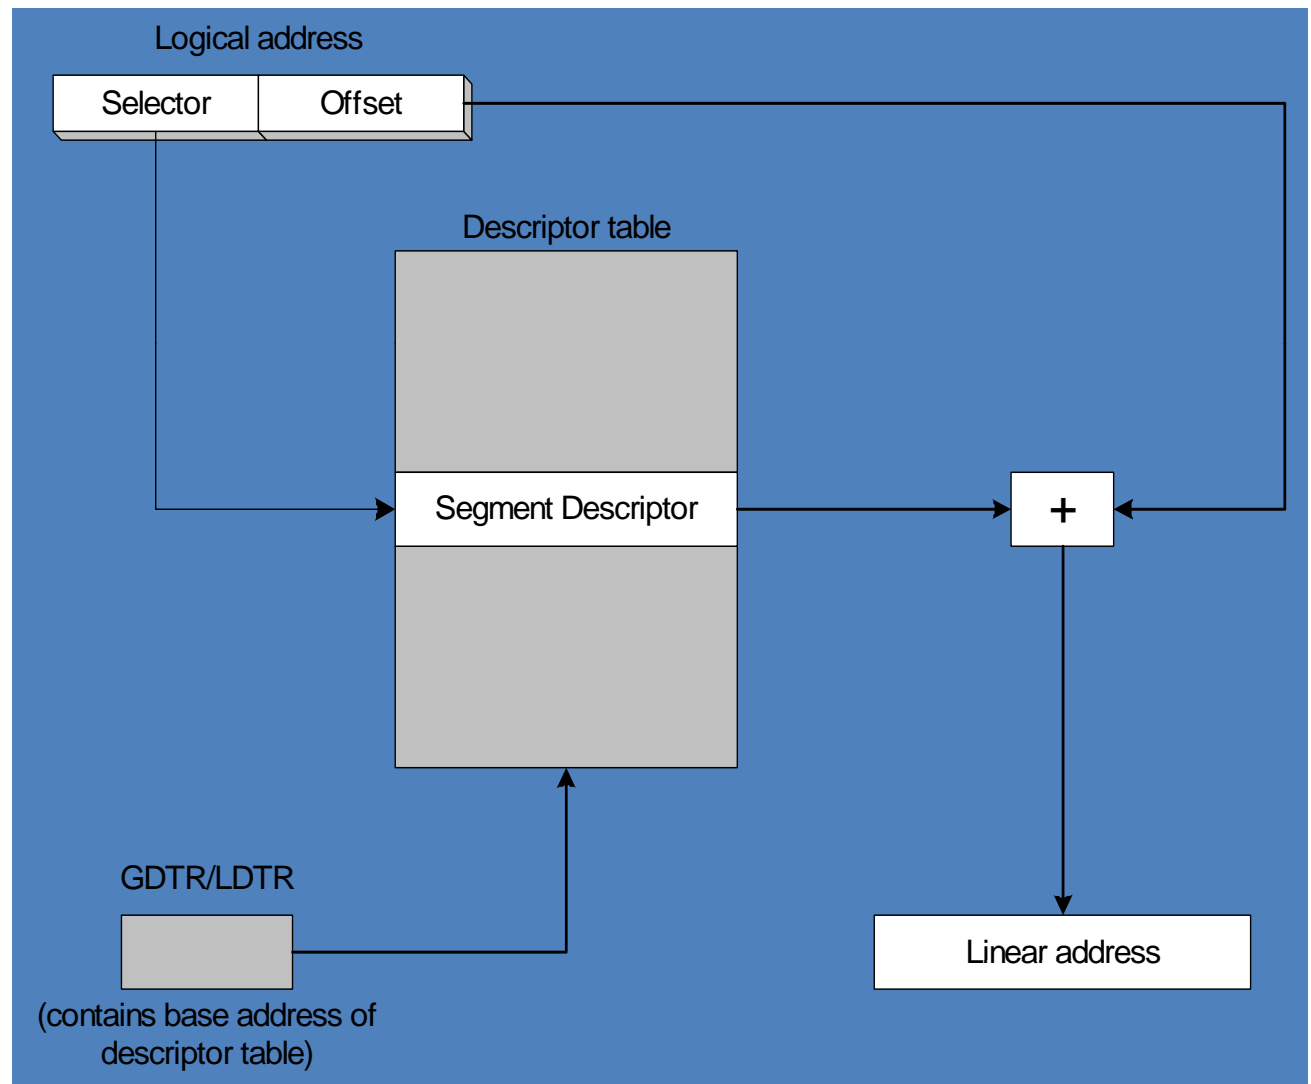


Translating Addresses (See Ch11.4)

- The IA-32 processor uses a one- or two-step process to convert a variable's logical address into a unique memory location.
- The first step combines a segment value (16-bit, segment register) with a variable's offset (32-bit) to create a **linear address**.
- The second optional step, called **page translation**, converts a linear address to a physical address.

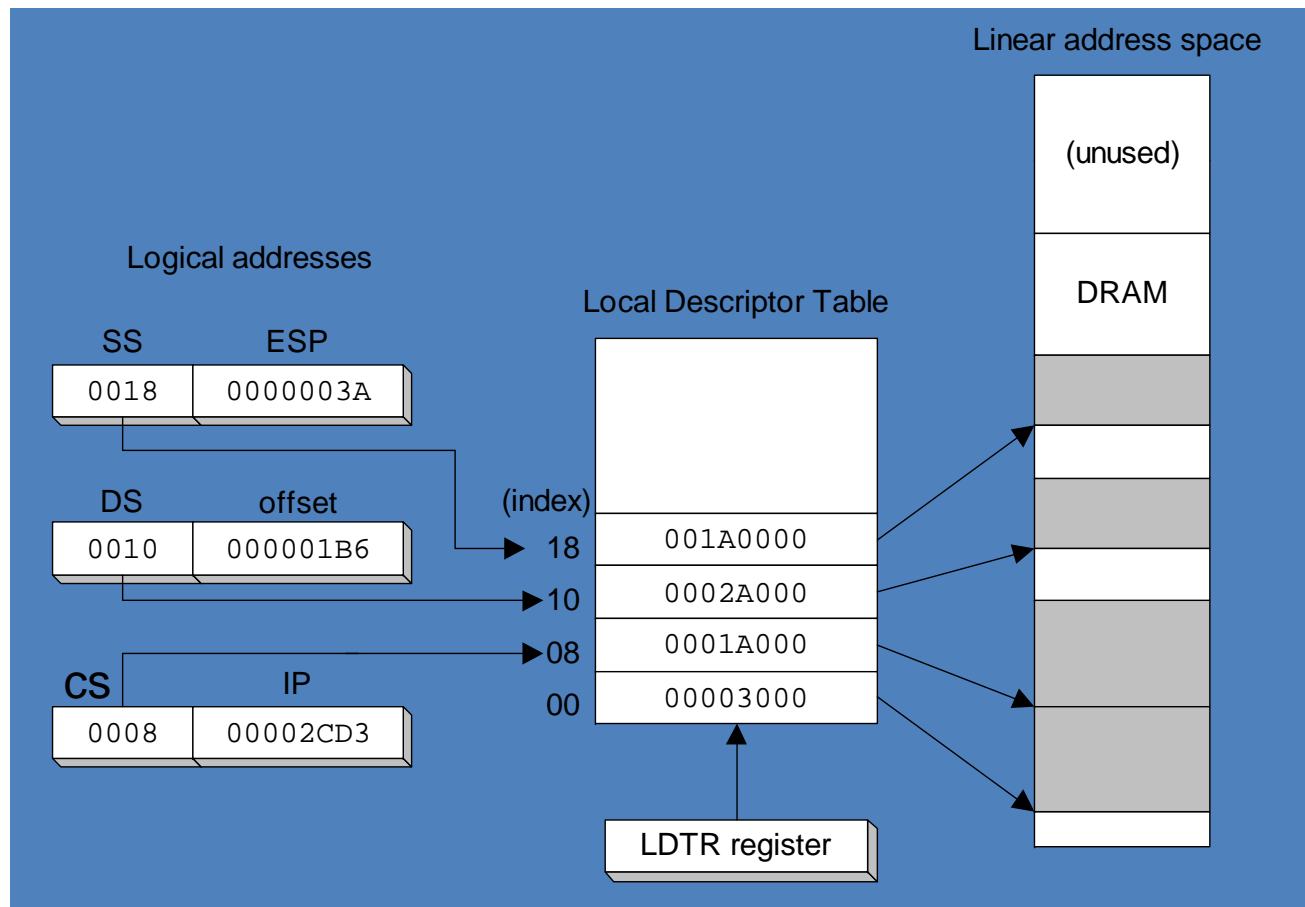
Converting Logical to Linear Address

The segment selector (16-bit) points to a **segment descriptor**, which contains the base address of a memory segment. The 32-bit offset from the logical address is added to the segment's base address, generating a 32-bit **linear address**.



Indexing into a Descriptor Table

- Each segment descriptor indexes into the program's local descriptor table (LDT)
- Each table entry is mapped to a linear address:



Paging

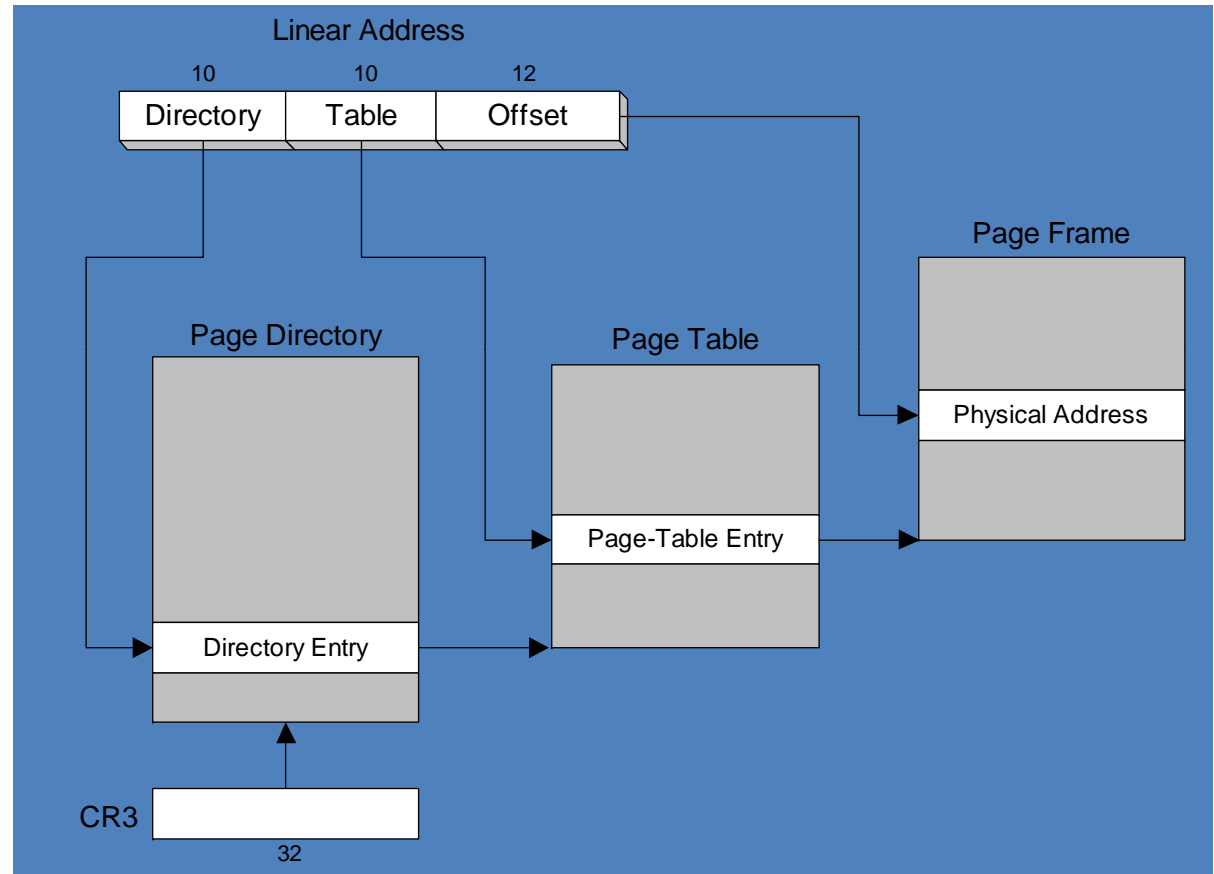
- Supported directly by the CPU
- Divides each segment into 4096-byte blocks called **pages**
- Sum of all programs can be larger than physical memory
- Part of running program is in memory, part is on disk
- **Virtual memory manager** (VMM) – OS utility that manages the loading and unloading of pages
- As the program runs, the processor selectively unloads inactive pages from memory and loads other pages that are immediately required.

Paging (cont.)

- OS maintains **page directory** and **page tables**
- **Page translation:** CPU converts the linear address into a physical address
- **Page fault:** issued by CPU when a needed page is not in memory, and CPU interrupts the program
- OS copies the page into memory, program resumes execution

Page Translation

A linear address is divided into a page directory field, page table field, and page frame offset. The CPU uses all three to calculate the physical address.



Paging

- Supported directly by the CPU
- Divides each segment into 4096-byte blocks called **pages**
- Sum of all programs can be larger than physical memory
- Part of running program is in memory, part is on disk
- **Virtual memory manager** (VMM) – OS utility that manages the loading and unloading of pages
- **Page fault** – issued by CPU when a page must be loaded from disk

Intel Microprocessor History

- Intel 8086, 80286
- IA-32 processor family
- P6 processor family
- CISC and RISC

Early Intel Microprocessors

- Intel 8080
 - 64K addressable RAM
 - 8-bit registers
 - CP/M operating system
 - S-100 BUS architecture
 - 8-inch floppy disks!
- Intel 8086/8088
 - IBM-PC Used 8088
 - 1 MB addressable RAM
 - 16-bit registers
 - 16-bit data bus (8-bit for 8088)
 - separate floating-point unit (8087)

The IBM-AT

- Intel 80286
 - 16 MB addressable RAM
 - Protected memory
 - several times faster than 8086
 - introduced IDE bus architecture
 - 80287 floating point unit

Intel IA-32 Family

- Intel386
 - 4 GB addressable RAM, 32-bit registers, paging (virtual memory)
- Intel486
 - instruction pipelining
- Pentium
 - superscalar, 32-bit address bus, 64-bit internal data path

Intel P6 Family

- Pentium Pro
 - advanced optimization techniques in microcode
- Pentium II
 - MMX (multimedia) instruction set
- Pentium III
 - SIMD (streaming extensions) instructions
- Pentium 4 and Xeon
 - Intel NetBurst micro-architecture, tuned for multimedia